

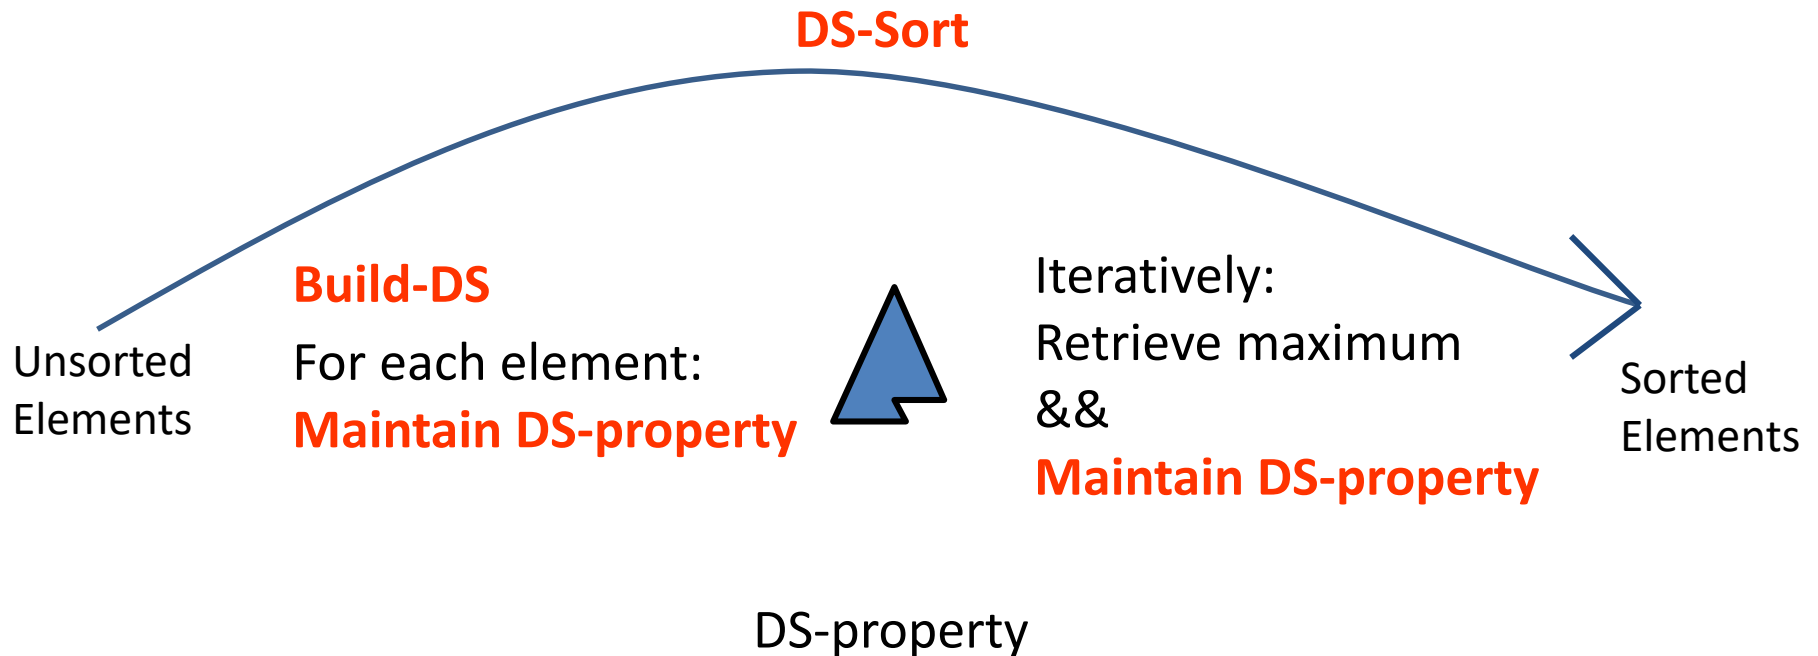
Outline

- Motivation: efficient sorting
- Trees
 - Binary Trees
- Heap
- Heapsort



HeapSort (One Slide)

using a data-structure (**DS**) as part of the algorithm design



(Recall) The Sorting Problem

Input: sequence of n numbers $\langle a_1, a_2, \dots, a_n \rangle$

Output: permutation (reordering) of the input
 $\sigma(a_i) = b_j$ such that $b_1 \leq b_2 \leq \dots \leq b_n$

Note: the sets $\{a_1, a_2, \dots, a_n\} = \{b_1, b_2, \dots, b_n\}$

Example:

- input: $\langle 31, 41, 59, 26, 41, 58 \rangle$
- output: $\langle 26, 31, 41, 41, 58, 59 \rangle$



Motivation For Sorting

- Among the most frequently used algorithms in CS
- Allows:
 - Binary search in $O(\log N)$ time
 - $O(1)$ time access to k^{th} largest element
 - Easy detection of any duplicates

Motivation For Heap Sorting

How to solve the sorting problem efficiently in terms of

- Time complexity
- Space complexity

Insertion-sort uses only constant number of extra storing cells (**But what is the worst case running time?**)

Merge-sort * takes $O(n \log n)$ but it is not *in-place*


We are seeking an efficient ADT that both operates *in-place* and takes $O(n \log n)$ time !

*coming soon

Trees

A **Tree** is a data structure with the following properties:

- Consisting of layers
- Has a single element in the top layer
- Each element in layer i points to elements in layer $i+1$
- Only a single element in layer i points to an element in layer $i+1$



How many edges in a tree with N nodes?

$N-1$ edges

Tree

Recursive Definition

A tree is a set of nodes, either

- It is an empty set of nodes, or
- It has one node called the **root** from which zero or more trees (subtrees) descend

Tree Terms

Child: B is a child of A iff A points to B

Parent: A is a parent of B iff A points to B

Sibling: A and B are siblings if they have the same parent

Root: a node with no parent

Leaf: a node with no children

Path: sequence of connected nodes

Examples of Tree Terms

A is the **root**

A-F are **nodes**

AD is an **edge** (one out of 5)

B,E,F,D are the **leaves**


A is the **parent** of C

E,F are **children** of C

B,C,D are **siblings**

A-C-E is a **path**

C is the root of the **subtree** consisting of C,E,F



More Tree Terms


Length of a path = number of edges between two nodes

Depth of a node A = length of path from root to A

Depth of tree = depth of deepest node

Height of node A = length of longest path from A to a leaf

Height of tree = height of the root



Implementation of Trees

Each node includes:

Option 1: a value + one pointer to each child


How many pointers should we allocate space for?

Other options?

Binary Tree




Each node has at most two children

- Popular data structure in computer science
- Will talk about it again latter in the course



Implementation of Binary Trees

Each node is implemented by a structure with value (key) and two pointers (left child & right child)



Binary Tree Trivia


Amongst all binary trees with N nodes:

- Which tree has the maximal depth?
- Which tree has the minimal depth?

Tree With Maximal Depth

Which tree has the maximal depth?


- Degenerate case: a linked list
- Depth = $N-1$



The Tree With Minimal Depth


Which tree has the minimal depth?

Is it unique?





Complete Binary Tree

Complete Binary Tree: All nodes are in use



Nearly Complete Binary Tree

(Nearly) Complete Binary Tree: All nodes are in use (except for possibly the right part of the bottom row.)



* Nearly Complete

Nearly Complete Binary Tree Trivia


Consider a nearly complete binary tree with N nodes:

- How many nodes in depth i ?
- How many nodes in height j ?
- What is the height(=depth) of the tree?

Nodes No. In Depth i

Consider a nearly complete binary trees with N nodes:

- How many nodes in depth i ?



Nodes No. In Height j : Intuition

Consider a nearly-complete binary tree with N nodes:

- How many nodes in total in a tree with depth d ?

$$a\left(\frac{r^n - 1}{r - 1}\right)$$

$$N = \sum_{i=0}^d 2^i = 2^{d+1} - 1$$

→ adding a complete level “~ doubles” the number of elements


Equivalently: How many elements in the deepest (it is a complete) level $h=0$?

Answer: approximately half

Nodes No. In Height j : Intuition

Consider a nearly complete binary trees with N nodes:

- How many nodes in height j ?




Formally: Height j includes at most $\left\lceil \frac{N}{2^{j+1}} \right\rceil$

Tree Height : Intuition

Consider a nearly complete binary trees with N nodes:

- What is the height(=depth) of the tree?



Tree Depth Analysis: Formally


Using the above observation:

- At depth i , there are 2^i nodes
- At depth d (tree depth), there may be 1 to 2^d nodes

Let N denote the total number of nodes:

$$\sum_{i=0}^{d-1} 2^i + 1 \leq N \leq \sum_{i=0}^{d-1} 2^i + 2^d$$

$$2^d \leq N \leq 2^{d+1} - 1$$



Depth Analysis

$$2^d \leq N \leq 2^{d+1} - 1$$


- From the left inequality: $d \leq \log N$
- From the right inequality: $\log(N + 1) \leq d + 1$
and: $\log(N) < d + 1$

In total: $\log(N) - 1 < d \leq \log N$

Equivalently: $d = \lfloor \log_2(N) \rfloor$

Nearly Complete Binary Trees

- How many nodes in depth d ? $\sim 2^d$
- How many nodes in height h ? $\frac{N}{2}, \frac{N}{2^2}, \dots$ $\sim \left\lceil \frac{N}{2^{h+1}} \right\rceil$
- What is the height(=depth) of the tree? $\sim \lfloor \log(N) \rfloor$




Heap

It is convenient to view a (binary) heap as a nearly complete binary tree.

The Max-heap property:


the key of the parent is equal or greater than the key of the children.



Heap Properties


- Binary heaps provide limited ordering information
- Each *path* is sorted, but *siblings are not sorted*
- Binary heap is \neq binary search tree (future ...)

This is a binary heap




Examples


Good




complete tree,
heap order is "max"



complete tree,
heap order is "min"



complete tree, but min
heap order is broken



not complete

Bad

Min / Max Heap

We focus on Max-heaps.

By symmetry, the statements, procedures and definitions are relevant for Min-heaps.

ADT: Max-Heap

Operations:

Empty (T)

Insert (T,x)


Min(Q)

Del_Max(Q)


Max-heap property: the key of the parent is equal or greater than the keys of both children

Heap Implementation

Binary Heap as array object:



The rules (of how to insert and delete elements) allow us to view it as a nearly complete binary tree.



Heap as a Complete Binary Tree

- Basic operations run in time that is proportional to the root height*
- The height is $\Theta(\log N)$ (as proved above):



The Heap ADT is potentially useful for sorting in order $O(N \log N)$ instead of $O(N^2)$

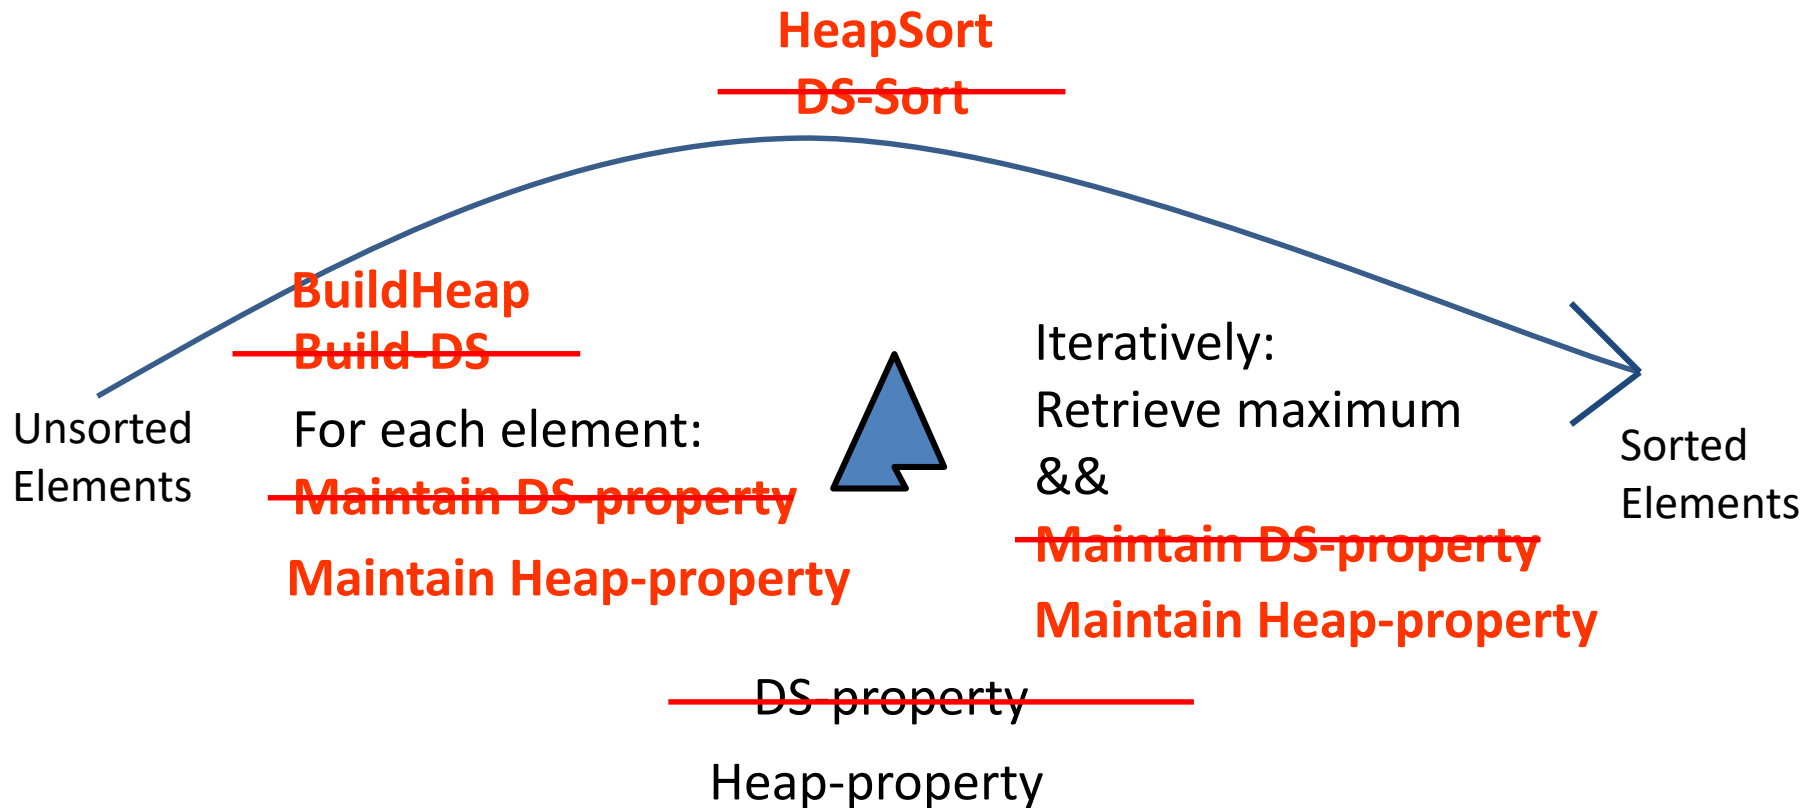
* We will prove or demonstrate it later in this lecture

Binary Heap Space Analysis

- Space needed for heap of at most $MaxN$ nodes: $O(MaxN)$
 - An array of size $MaxN$, plus a variable to store the current size N of the heap



HeapSort (One Slide)

Heap
using a data-structure (~~DS~~) as part of the algorithm design



Max-Heapify Pseudocode


MAX-HEAPIFY(A, i)

1	$l = \text{LEFT}(i)$		Largest = Index of the node with highest values
2	$r = \text{RIGHT}(i)$		
3	if $l \leq A.\text{heap-size}$ and $A[l] > A[i]$		
4	$largest = l$		
5	else $largest = i$		
6	if $r \leq A.\text{heap-size}$ and $A[r] > A[largest]$		Percolate into deeper level
7	$largest = r$		
8	if $largest \neq i$		
9	exchange $A[i]$ with $A[largest]$		
10	MAX-HEAPIFY($A, largest$)		


Cormen 6.2

Max-Heapify Visualization


Running time of Max-Heapify on a node of height h is $O(h)$



Max-Heapify (A,2)



Max-Heapify (A,4)



Max-Heapify (A,9)

Cormen 6.2

Build-Heap Pseudocode

BUILD-MAX-HEAP(A)






```
1   $A.heap-size = A.length$   
2  for  $i = \lfloor A.length/2 \rfloor$  downto 1  
3      MAX-HEAPIFY( $A, i$ )
```




$\lfloor \frac{N}{2} \rfloor$ is the index of the last non-leaf

Build Heap Visualization

Idea: using Max-Heapify in a bottom-up manner
from $\left\lfloor \frac{N}{2} \right\rfloor$ to 1 (e.g. from 5 to 1, when $N=10$)



Result: Heap



Cormen 6.3

Correctness of Build Heap

- For $n=1$, a tree with a single node is a heap.

Hence, trees rooted at $\left\lfloor \frac{n}{2} \right\rfloor < i \leq n$ are heaps.

- For $i \leq \left\lfloor \frac{n}{2} \right\rfloor$, children are heaps.

Hence, after calling Max-Heapify we obtain heap.

Time Complexity Build Heap

Claim: The running time of Build-Max-Heap is $O(n)$

Proof :

BuildHeap performs **Heapify** on nodes with height $j \geq 1$
Heapify on node of height j costs at most cj

Recall:

- at most $\left\lceil \frac{N}{2^{j+1}} \right\rceil$ elements at level in height j
- the height of the tree is at most $\lfloor \log(N) \rfloor$

$$\sum_{j=1}^{\lfloor \log(n) \rfloor} \left\lceil \frac{N}{2^{j+1}} \right\rceil cj$$

Time Complexity of Build Heap

$$\sum_{j=1}^{\lfloor \log(N) \rfloor} \left\lfloor \frac{N}{2^{j+1}} \right\rfloor c_j \leq$$

$$\sum_{j=0}^{\lfloor \log(N) \rfloor} \frac{N}{2^j} c_j = cN \sum_{j=0}^{\lfloor \log(N) \rfloor} \frac{j}{2^j}$$

$$< cN \sum_{j=0}^{\infty} \frac{j}{2^j} = 2cN$$

$$\sum_{j=0}^{\infty} \frac{j}{2^j} = 2$$

Heapsort Pseudocode

HEAPSORT(A)

1 BUILD-MAX-HEAP(A)

2 for $i = A.length$ downto 2

3 exchange $A[1]$ with $A[i]$

4 $A.heap-size = A.heap-size - 1$

5 MAX-HEAPIFY($A, 1$)

Iteratively:

Retrieve maximum

&&

Maintain DS-property


Retrieve Maximum

Retrieve maximum: Easy!

- Return root value (that is $A[1]$)
- Run time = ?

Maintain DS-property: harder!

- why ?




The Challenge of Retrieve Maximum

Challenge: Following retrieve maximum (remove root), the tree will have one less node → temporarily it is not a tree

Solution: Replace the root with the “last element” yet not sorted


Why is it not a Heap?



Iteratively:
Retrieve maximum
&&

Maintain DS-property

Heapsort Visualizaion



Recommendation: Cormen 6.4

Time Complexity HeapSort

Claim: The running time of HeapSort is $O(n \log n)$

Proof :

- **HeapSort** performs **Max-Heapify** on $n-1$ nodes
- Each **Max-Heapify** costs at most $c \log n$

In total, **HeapSort** costs at most $cn \log n$.

Summary – What Is Heap?

The Heapsort Idea: using a data structure that

- returns the maximum in $O(1)$
- maintains the heap property in $O(\log n)$.

Heap ~ Nearly Complete Tree

Max-heap property: the key of the parent is equal or greater than the keys of both children.

Summary – HeapSort Analysis

using a data-structure(**DS**) as part of the algorithm design

